

## FFT's

My primary interest outside of computers is music and I've invested/wasted a significant amount of time over the course of my life using DAW's (digital audio workstations) and other audio processing tools. This sent me searching for an algorithm fundamental to the processing and quickly settling on the Fast Fourier Transform or FFT for short. Like all transforms the purpose of a Fourier transform is to map a function from its original function space to another function space. In the case of a Fourier Transform the goal is to take a function and decompose it into its constituent sinusoidal frequencies. Since any waveform can be reproduced by adding sinusoidal functions together this can be useful for processing and making sense of complex functions and waveforms. A simple example that I will be using throughout this project is interpreting the waveform(signal or amplitude over time) of a musical signal or chord and expressing it instead in terms of the volumes and frequencies of its constituent notes so that a filter can be applied to attenuate or remove one of those constituent inputs.

A project folder for a single song may have Gigabytes of raw audio information that has to get processed before it gets rendered down into a familiar and portable format like mp3. Making optimal use of the hardware available to consumers so that work can be performed as quickly as possible is an important job of any standard daw. On today's computer systems that means making use of the multiple cores available to try and speed up an algorithm that may get asked to perform on very large data sets. In our case, this means FFT's.

An FFT is an algorithm that computes the Discrete Fourier Transform of a sequence in  $O(N \log N)$  instead of  $O(N^2)$  as would occur when one simply applies the definition of DFT. Our goal is to further improve this through parallelization to achieve the smallest time complexity

possible, which should be  $O(n/p \log p)$  where  $p$  is our number for threads.

The most common FFT algorithm in use today is known as the Cooley-Tukey algorithm of which there are many varieties. Some of the simpler easier to grasp, less space-efficient varieties provide less opportunity for speed up via parallelization. The simplest out of place recursive version is known as the radix-2 DIT case doesn't present much if any opportunity for parallelization at all. However, there are a number of iterative approaches that attempt to limit the space complexity to in-place that make use of multiple nested loops and has several of the steps decomposed into separate loops instead of all occurring in one. These present ample opportunity to see if parallelization attempts can be made. I will be attempting to parallelize an Iterative radix-2 FFT that is implemented using bit reversal permutation based on this pseudocode

**algorithm** iterative-fft **is**

**input:** Array  $a$  of  $n$  complex values where  $n$  is a power of 2.

**output:** Array  $A$  the DFT of  $a$ .

bit-reverse-copy( $a$ ,  $A$ )

$n \leftarrow a.length$

**for**  $s = 1$  **to**  $\log(n)$  **do**

$m \leftarrow 2^s$

$\omega_m \leftarrow \exp(-2\pi i/m)$

**for**  $k = 0$  **to**  $n-1$  **by**  $m$  **do**

$\omega \leftarrow 1$

**for**  $j = 0$  **to**  $m/2 - 1$  **do**

$t \leftarrow \omega A[k + j + m/2]$

$u \leftarrow A[k + j]$

$A[k + j] \leftarrow u + t$

$A[k + j + m/2] \leftarrow u - t$

$\omega \leftarrow \omega \omega_m$

```
return A
```

**algorithm** bit-reverse-copy( $a, A$ ) **is**

**input:** Array  $a$  of  $n$  complex values where  $n$  is a power of 2.

**output:** Array  $A$  of size  $n$ .

```
 $n \leftarrow a.length$ 
```

```
for  $k = 0$  to  $n - 1$  do
```

```
     $A[\text{rev}(k)] := a[k]$ 
```

While Fourier Transforms are calculated through integration the process behind rapidly executing the discrete Fourier transform involves factorization using matrix multiplication. This pattern of matrix multiplication when drawn out as a diagram creates a pattern that looks like a butterfly which is why it's often referred to as a series of butterfly operations in the literature. During these butterfly equations, our inputs are multiplied by a series of multiplicative trigonometric constants that are referred to in the literature as "Twiddle Factors". One of the ways we will be attempting to induce speed up is by separating the calculation of these twiddle factors into a table before run so the process can be effectively parallelized. Many simplified versions of the code calculate these factors during each iteration of the main loop and thus become a hurdle to parallelization this way.

## Sequential Solution and Computation Time

```
jkl@funRuiner:~/Documents/final_proj/src$ make
g++ -g -Wall -o serial_fft serial_fft.cpp
g++ -g -Wall -fopenmp -o omp_fft omp_fft.cpp -lm
jkl@funRuiner:~/Documents/final_proj/src$
```




240 seconds( 4min song )	1 Thread
	19.506
	19.513
	18.23
	19.216
	17.782
	18.874
	18.261
	19.554
<b>AVG OF 8 RUNS</b>	<b>18.867 seconds</b>

[illegible][illegible]

# Speed Up: A Closer Look

S(n) = 1 / ((1-P) + P/n)									
Action Time		Actual Speed up		85% Assumed		90% Assumed		95% Assumed	
1	18.867	1	1	1	1	1	1	1	1
2	10.133	1.861	1.74	1.82	1.9				
4	5.245	3.597	2.76	3.06	3.48				
8	3.14	6.008	3.9	4.71	5.93				
12	2.69	7.013	4.53	5.71	7.74				
16	2.382	7.92	4.92	6.4	9.14				
		8 Cores + SMT		8 Cores + SMT		8 Cores + SMT		8 Cores + SMT	
		12 Threads		16 Threads		20 Threads		24 Threads	
AMD SMT == INTEL HYPERTHREADING. First Generation AMD Ryzen Processors 1300/1500/1700 took a small but measurable single threaded hit with SMT enabled. Any time we push over 8 threads there is potentially risk for single thread performance degradation. Not sure if this will come into play but making notes of it.									
Notes on Data Selection									
After doing 8 consecutive runs per thread allotment I did another set of runs for outlier replacement. Selectively replacing any obvious high point outliers from each category. The idea is there is unlikely to be serious error of measurement in the negative direction but competing resources on my machine and the scheduler could skew the data high. I didn't bother doing a Standard deviation check, just intuitively removed obvious outliers.									
AVG OF 8 RUNS		18.867 seconds	10.133 seconds	5.245 seconds	3.1405 seconds	2.69 Seconds	2.382 Seconds		




I used a table to estimate parallelization against my real world data and kept adjusting P until I reached the best fit I could. There was no perfect P that fit all data points perfectly but 94% seems to be the sweet spot. It's notable performance curve worsens as thread count gets closer to the total cores available. There is still improvement but lets and less as we go further.

S(n) = 1/ ((1-P) + P/n)									
Action Time		Actual Speed up	85% Assumed	Amhdahls Law Speed UP at		90% Assumed	95% Assumed		
1	18.867	1	1	1	1	1	1	1	
2	10.133	1.861	1.74	1.82	1.9				
4	5.245	3.597	2.76	3.08	3.48				
8	3.14	6.008	3.9	4.71	5.93				
12	2.69	7.013	4.53	5.71	7.74				
16	2.382	7.92	4.92	6.4	9.14				
</									


[https://docs.google.com/spreadsheets/d/1\\_v3FaX-F1XovysnmPOWQ9fInSAZlukWGBXyd2aWtQG0/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1_v3FaX-F1XovysnmPOWQ9fInSAZlukWGBXyd2aWtQG0/edit?usp=sharing)

## Matching Each Point to an Efficiency Curve




## Comparative Speed Analysis.

My 94% parallelization vs 100%



My 94% parallelization vs theoretical best case scenario for FFTs of 96.7% according to <https://ilias.kuleuven.be/retrieve/200966>



In the end, we have 6% of the function that is serial and around 94% of it can be made parallel.

The program made maximum use of all my cores. The beginning and end of the program run it is single-threaded, simply generating signals and writing them down. The actual FFT portion you

can tell is running when all of your available cores spike in utilization together(or as many as you selected as an input parameter). Example derived by running **mpstat -P ALL 1**

08:16:03 PM	CPU	%usr	%nice	%sys	%iowait	%irq	%soft	%steal	%guest	%gnice	%idle
08:16:04 PM	all	96.07	0.00	0.25	0.00	0.00	0.00	0.00	0.00	0.00	3.68
08:16:04 PM	0	100.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
08:16:04 PM	1	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	2	95.05	0.00	0.99	0.00	0.00	0.00	0.00	0.00	0.00	3.96
08:16:04 PM	3	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	4	95.05	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.95
08:16:04 PM	5	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	6	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	7	96.04	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	3.96
08:16:04 PM	8	98.00	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00
08:16:04 PM	9	95.96	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.04
08:16:04 PM	10	95.96	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.04
08:16:04 PM	11	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	12	95.05	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.95
08:16:04 PM	13	96.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	4.00
08:16:04 PM	14	94.95	0.00	1.01	0.00	0.00	0.00	0.00	0.00	0.00	4.04
08:16:04 PM	15	95.05	0.00	0.99	0.00	0.00	0.00	0.00	0.00	0.00	3.96

The first few frames of mpstat will show only one core being utilized while that core generates the signal. Then when the FFT portion kicks in utilization spikes on all cores.

By shrinking the time to process a 4 min piece of “audio” from 20 seconds to 2 seconds plus some change we have not only produced a measurable difference in speed but one that could be felt any time an FFT needs to be performed.

### Spreadsheets used to track data and calculate speed up based on Amdahl’s Law

[https://docs.google.com/spreadsheets/d/1\\_v3FaX-F1XovysnmPOWQ9fInSAZlukWGBXyd2aWtQG0/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1_v3FaX-F1XovysnmPOWQ9fInSAZlukWGBXyd2aWtQG0/edit?usp=sharing)

### Slides that were used in my video and to construct this document.

[https://docs.google.com/presentation/d/1uNcFvTCgEbdALvfrYZXjPSW26iA1Jd4aeP6kOWofAo8/edit#slide=id.g107f9645bd5\\_0\\_15](https://docs.google.com/presentation/d/1uNcFvTCgEbdALvfrYZXjPSW26iA1Jd4aeP6kOWofAo8/edit#slide=id.g107f9645bd5_0_15)

### Youtube Link(please do not share )

<https://youtu.be/Asoy12T2egg>